

What is Cloud Computing?

Cloud computing is a model for enabling ubiquitous, convenient, **on-demand** network access to a shared pool of configurable computing resources (e.g., networks, servers, storage, applications, and services) that can be rapidly provisioned and released with minimal management effort or service provider interaction [1].

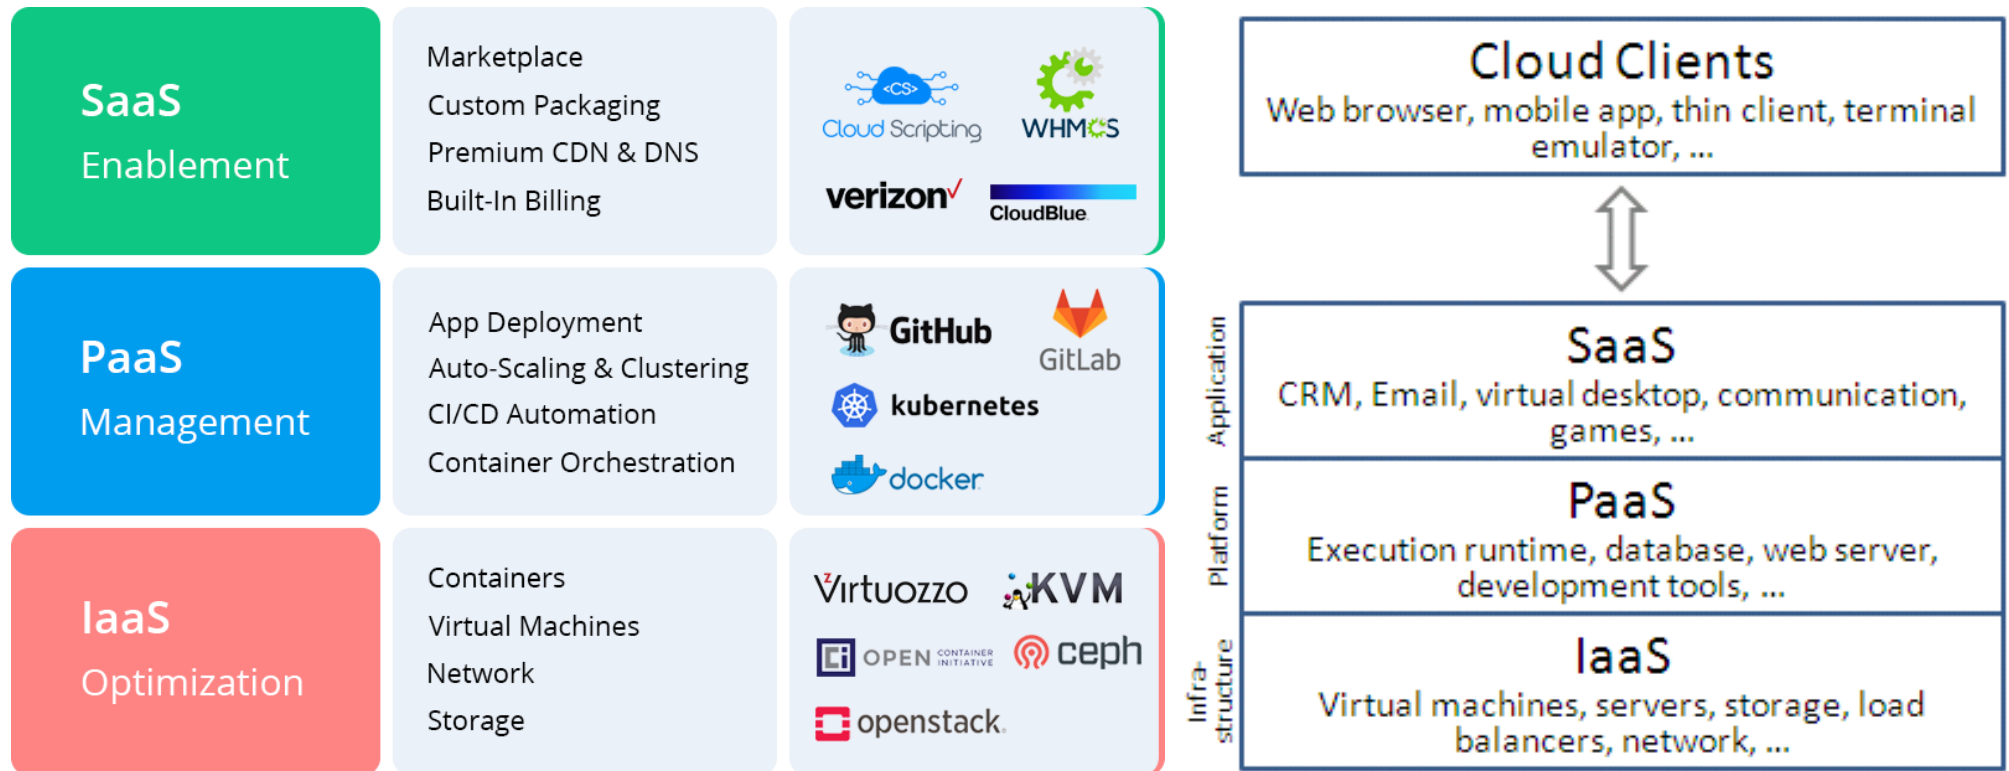
Informal: Cloud Computing = computing with large datacenters.

Our focus: Cloud Computing = computing as a utility outsourced to a third party or internal org

[1] Peter Mell (NIST), Tim Grance (NIST),
The National Institute of Standards and Technology (NIST) Definition of Cloud Computing
<https://csrc.nist.gov/publications/detail/sp/800-145/final>

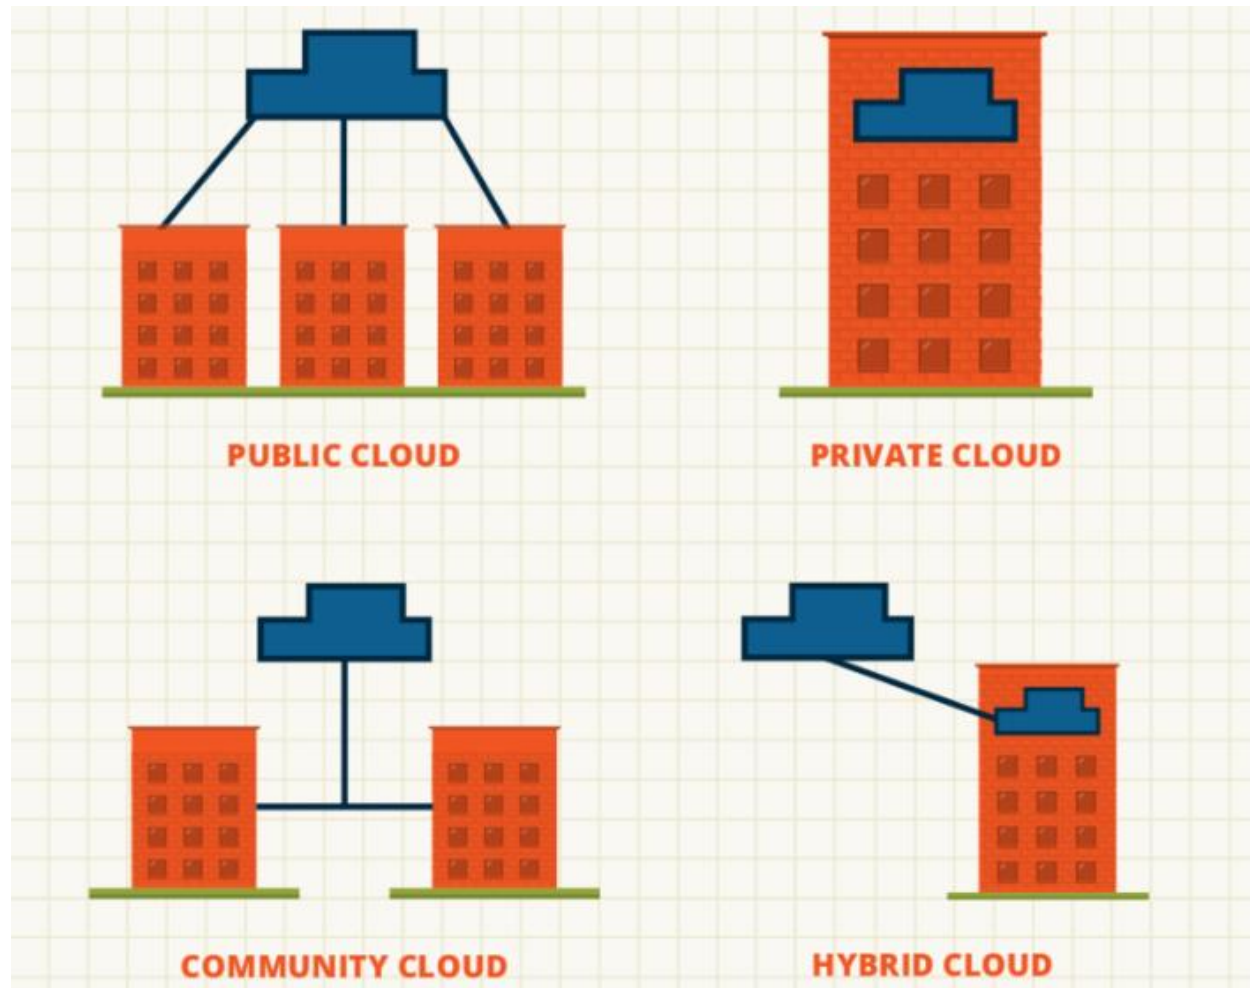
Service Models

Infrastructure as a Service (IaaS), Platform as a Service (PaaS), Software as a Service (SaaS)



Deployment Models


**Infrastructure as a Service (IaaS), Platform as a Service (PaaS),
Software as a Service (SaaS)**



Cloud Economics: For Users

Pay-as-you-go (usage-based) pricing:

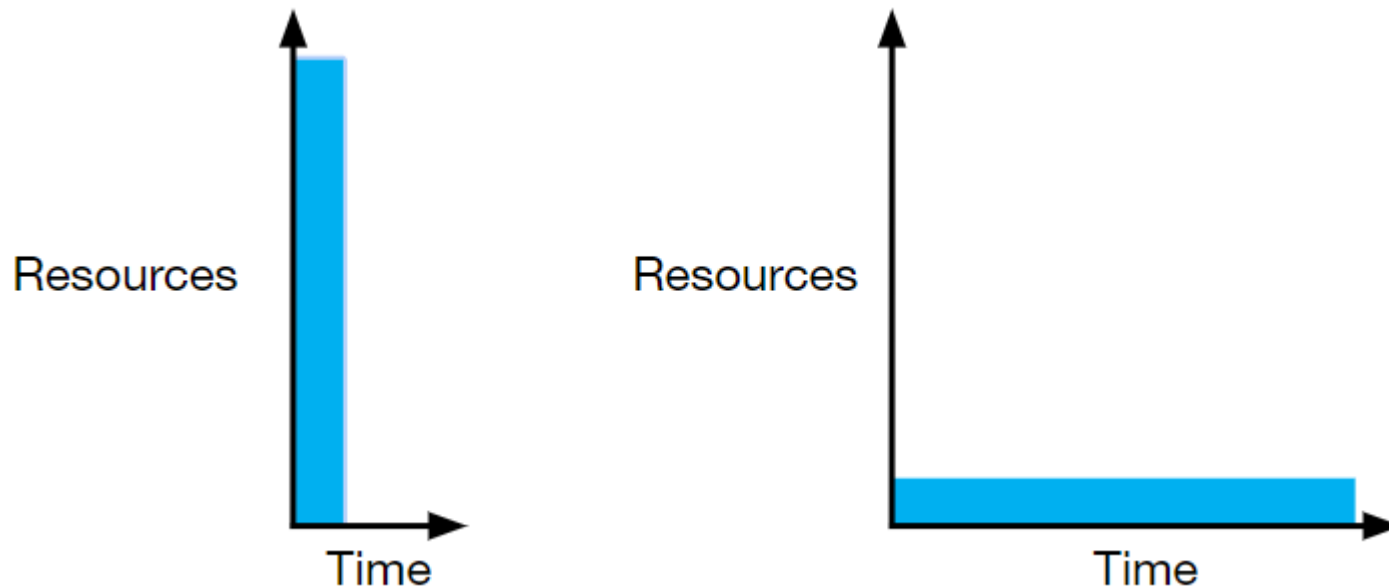
- Most services charge per minute, per byte, etc
- No minimum or up-front fee
- Helpful when apps have variable utilization



Cloud Economics: For Users

Elasticity:

- Using 1000 servers for 1 hour costs the same as 1 server for 1000 hours
- Same price to get a result faster!



Cloud Economics: For Providers

Economies of scale:

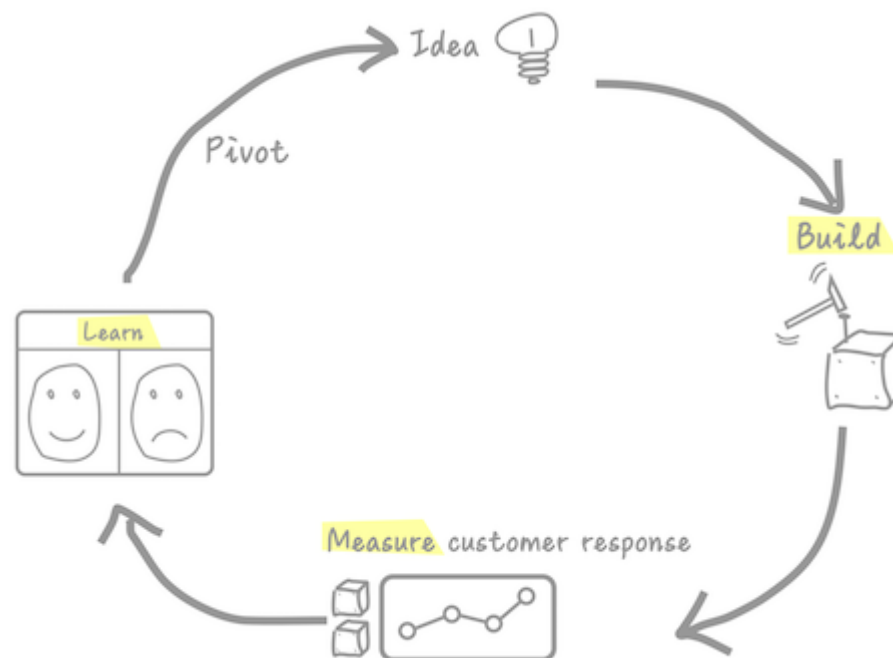
- Purchasing, powering & managing machines at scale gives lower per-unit costs than customers'
- To find Tradeoff: fast growth vs efficiency
- To find Tradeoff: flexibility vs cost



Cloud Economics: For Providers

Speed of iteration:

- Software as a service means fast time-to-market, updates, and detailed monitoring/feedback
- Compare to speed of iteration with ordinary software distribution



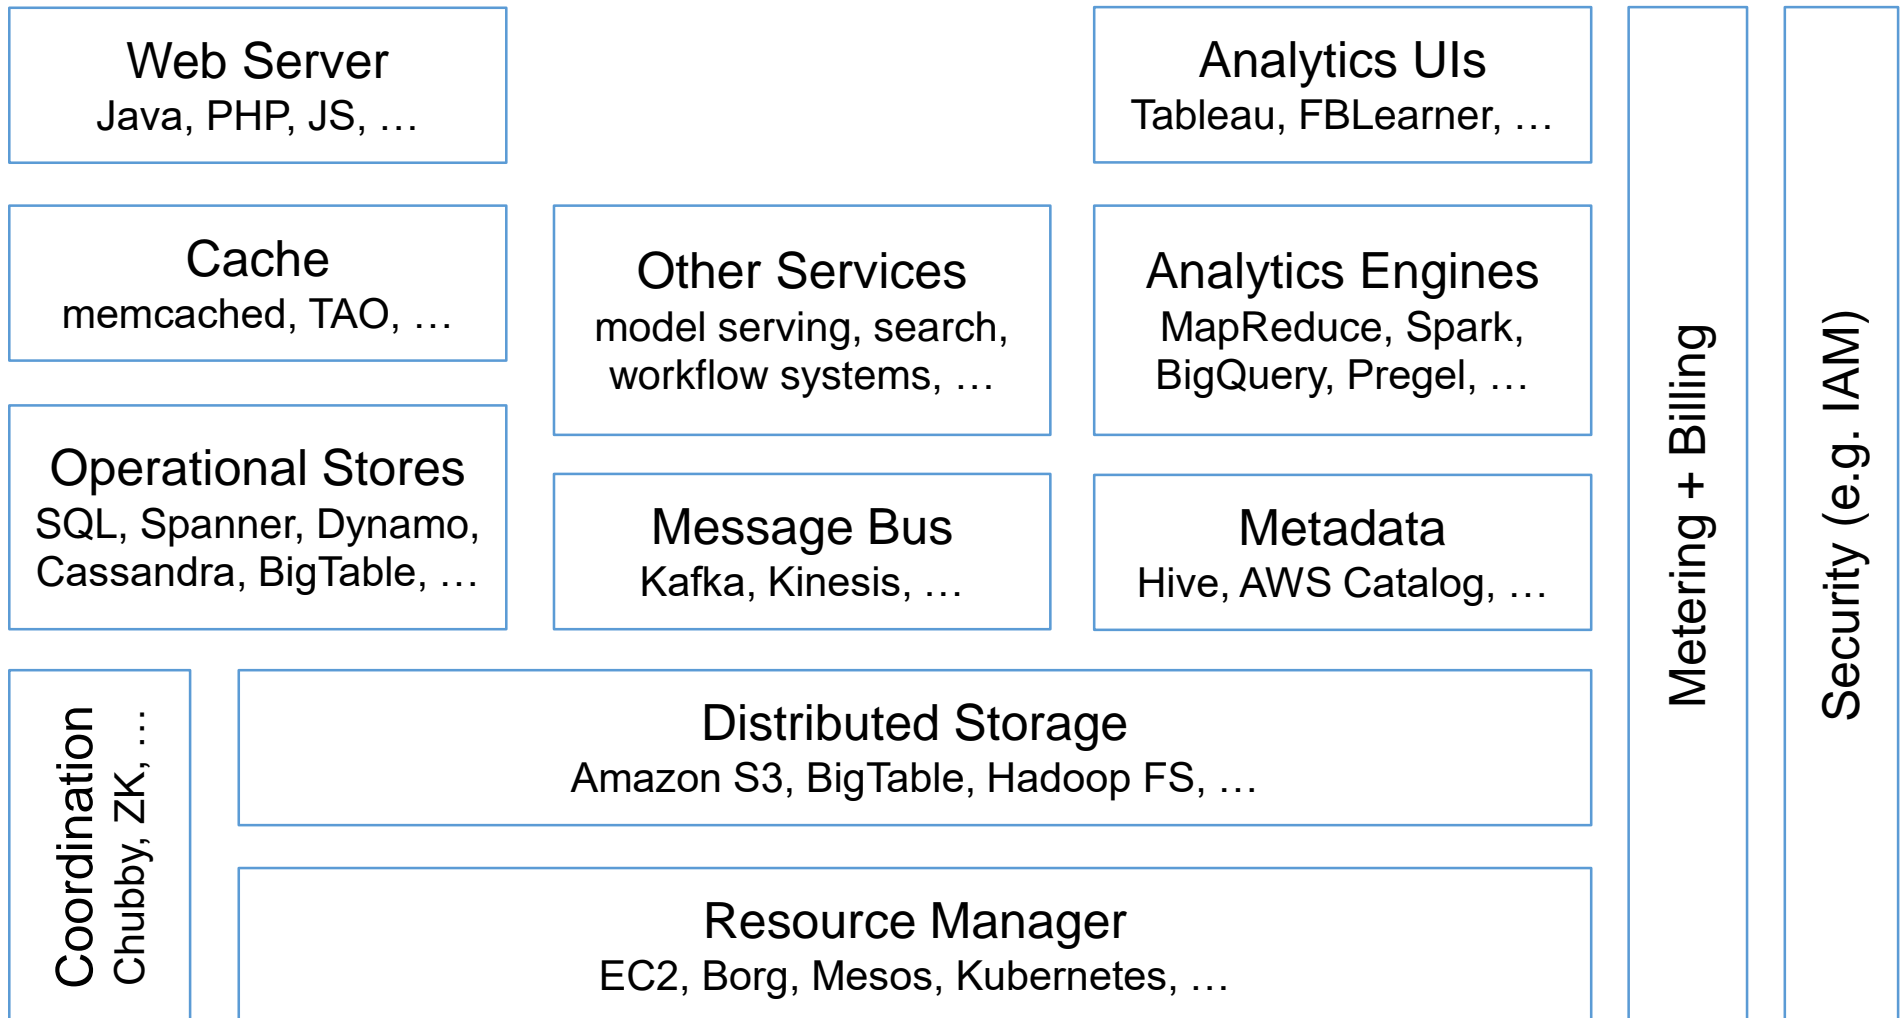
Other Interesting Features

- Spot market for preemptible machines, the ability to purchase a productive old server
- Wide geographic access for disaster recovery and speed of access
- Ability to quickly try exotic hardware
- Ability to A/B test anything

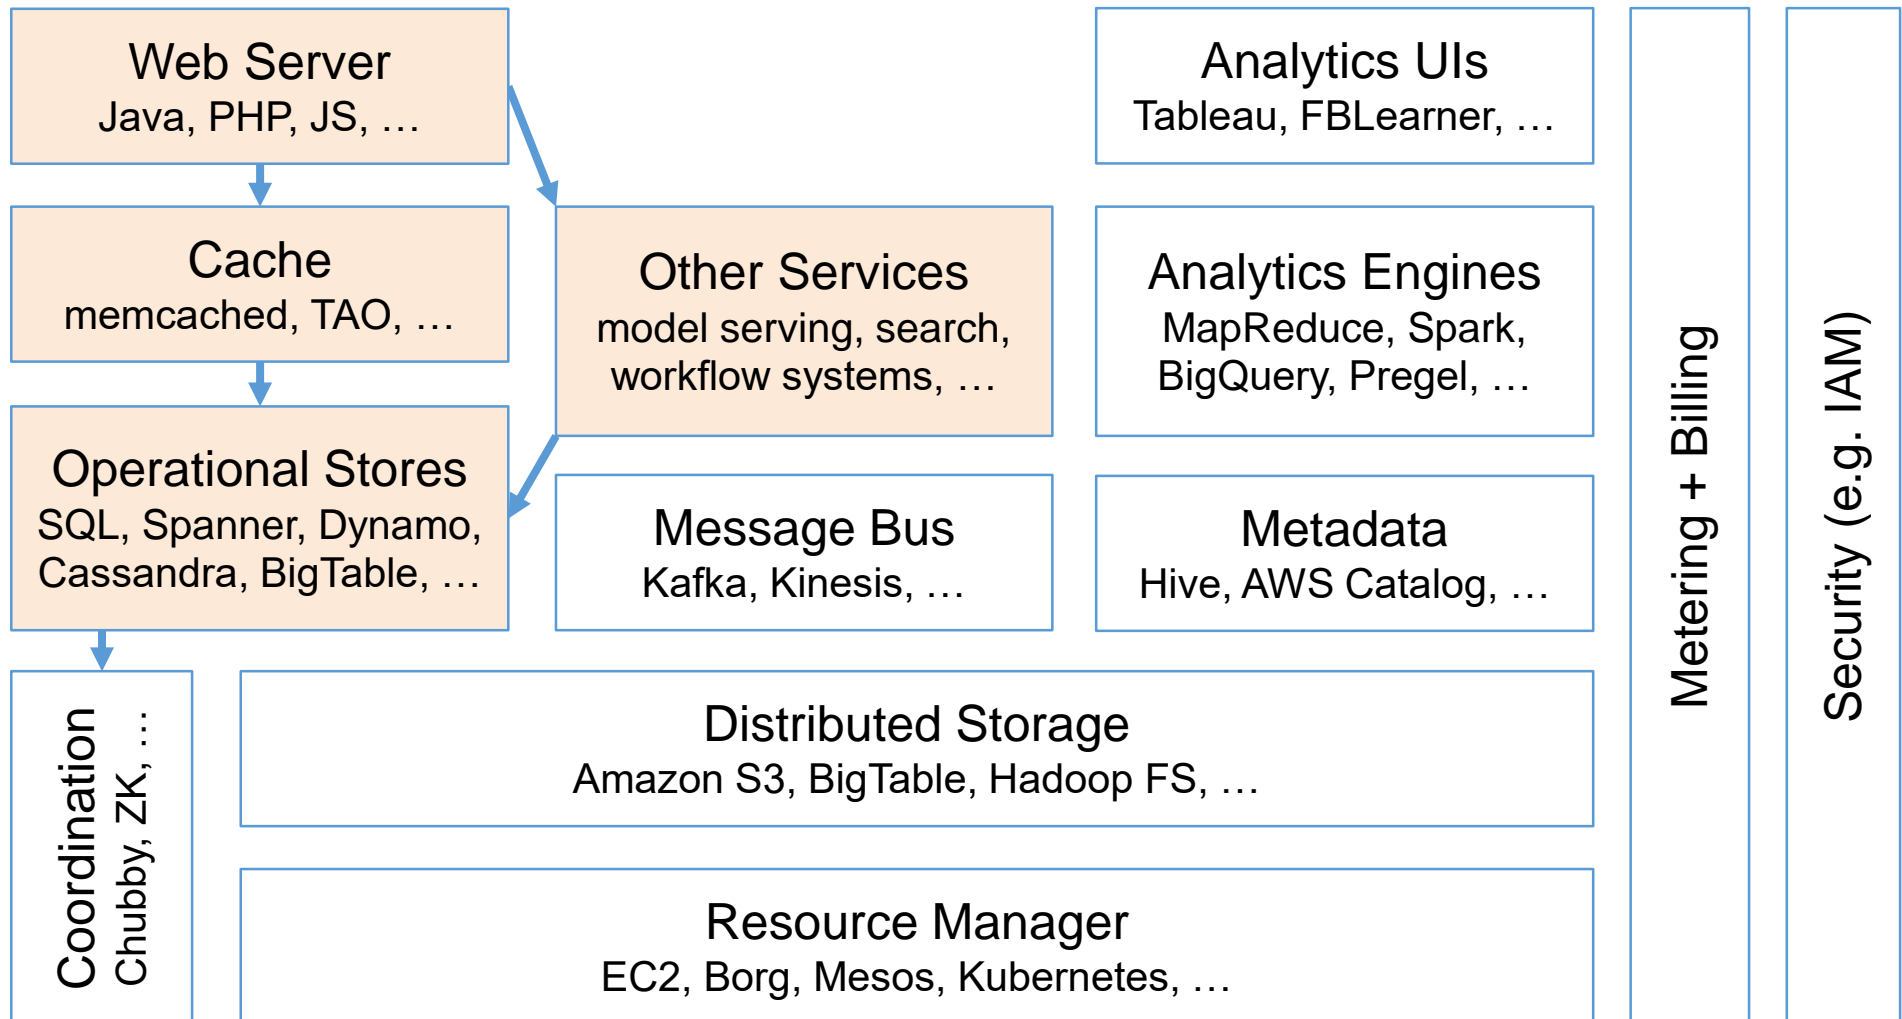
Common Cloud Applications

- Web and mobile applications
- Data analytics (MapReduce, SQL, ML, etc)
- Stream processing
- Batch computation (HPC, video, etc)

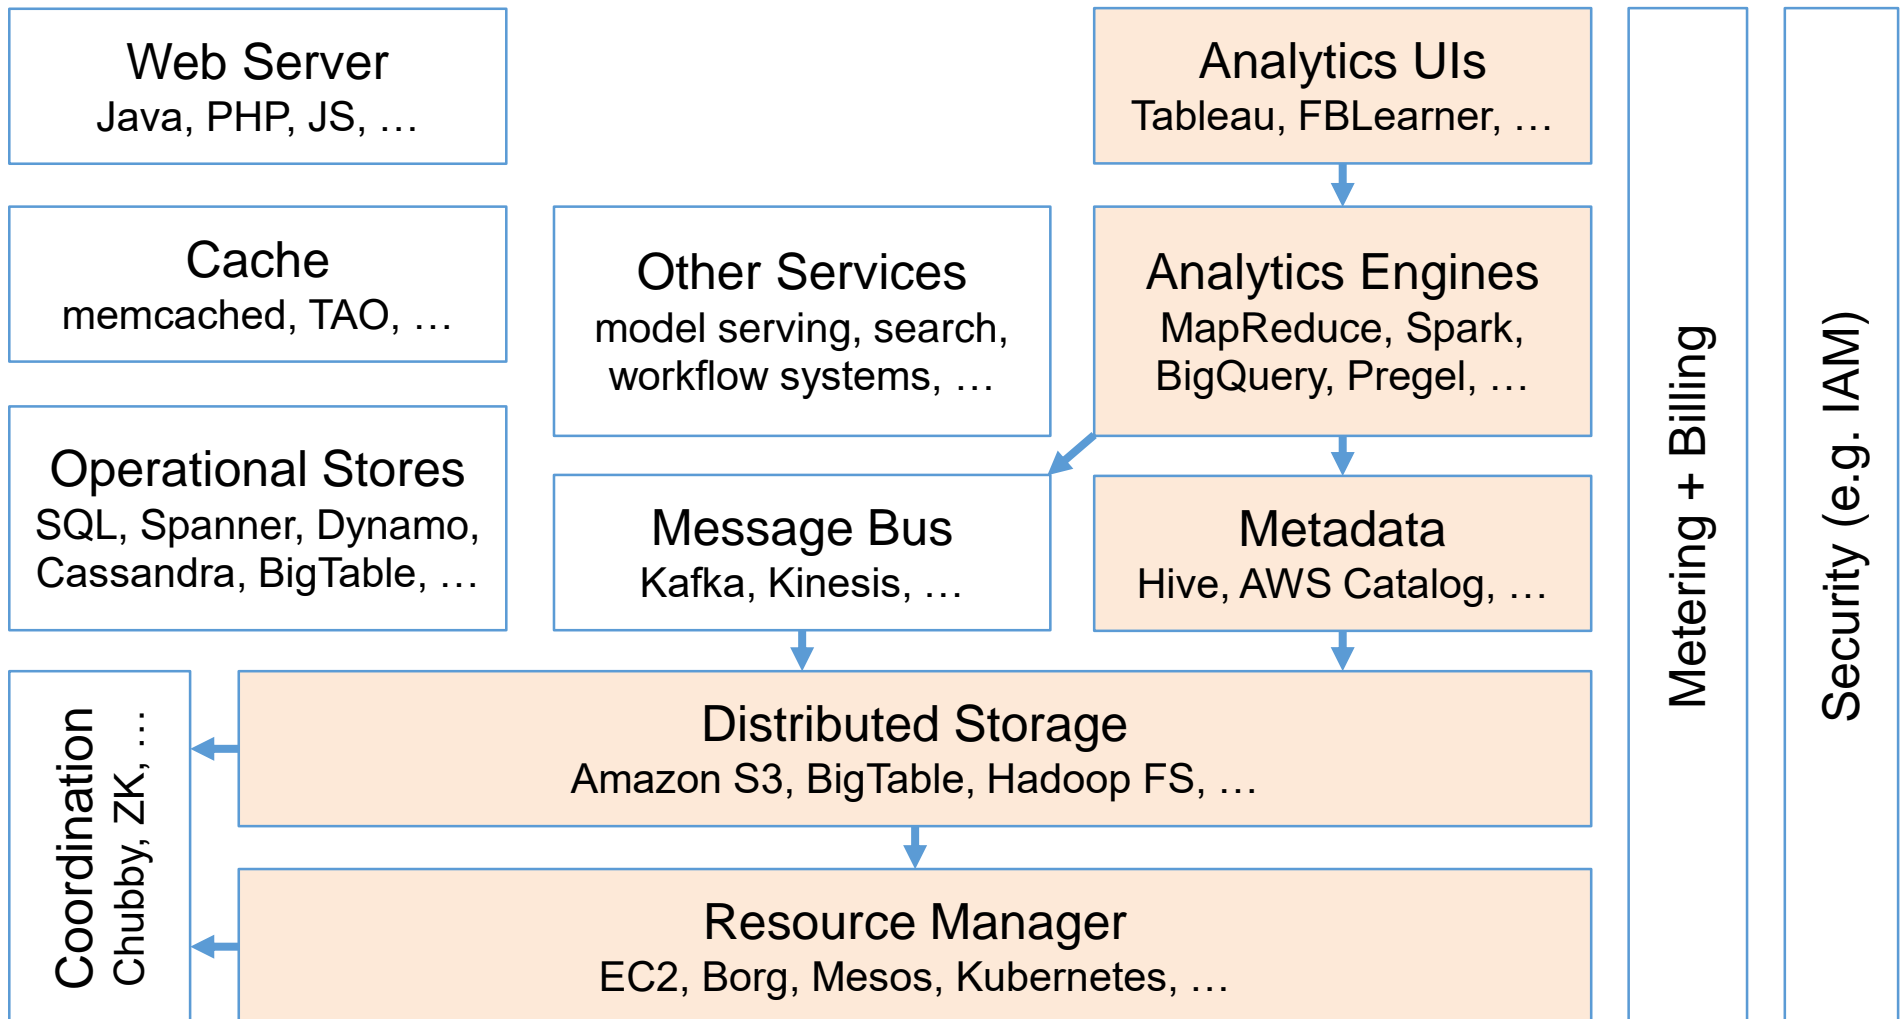
Cloud Software Stack




Example: Web Application



Example: Analytics Warehouse



Datacenter Hardware



Rows of rack-mounted servers

Datacenter: 50 – 200K of servers, 10 – 100MW

Often organized as few and mostly independent clusters

Datacenter Example



Datacenter HW: Compute

The basics

Multi-core CPU servers

1 & 2 sockets



What's new

GPUs

FPGAs

Custom accelerators (AI)

2-socket server



Hardware Heterogeneity

| Standard Systems | I Web | III Database | IV Hadoop | V Haystack | VI Feed |
|------------------|---------------------|---------------------------|-----------------------|-----------------------|-------------------------------------|
| CPU | High 2 x E5-2670 | High 2 x E5-2660 | High 2 x E5-2660 | Low 1 x E5-2660 | High 2 x E5-2660 |
| Memory | Low 16GB | High 144GB | Medium 64GB | Med-Hi 96GB | High 144GB |
| Disk | Low 250GB | High IOPS 3.2 TB Flash | High 15 x 4TB SATA | High 30 x 4TB SATA | Medium 2TB SATA + 1.6TB Flash |
| Services | Web, Chat | Database | Hadoop | Photos, Video | Multifeed, Search, Ads |

[Facebook server configurations]

Custom-design servers

Configurations optimized for major app classes

Few configurations to allow reuse across many apps

Roughly constant power budget per volume


Datacenter HW: Storage

The basics

Disk trays

SSD & NVM Flash

NVMe Flash



JBOD disk array




What's new

Non-volatile memories

New archival storage (e.g., glass)

NVM DIMM



Distributed with compute or NAS systems

Remote storage access for many use cases (why?)

Datacenter HW: Networking

The basics

10, 25, and 40GbE NICs

40 to 100GbE switches

Clos topologies

40GbE Switch



What's new

Software defined networking


Smart NICs

FPGAs

Smart NIC



Microsoft Catapult



Performance Metrics

Throughput

- Requests per second

- Concurrent users

- Gbytes/sec processed

- ...

Latency

- Execution time

- Per request latency

Useful Throughput Numbers

| | |
|---|---------------|
| DDR4 channel bandwidth (DDR4-2400 PC4-19200) | ~20 GB/sec |
| PCIe gen3 (2010) x16 channel | 15.8 GB/sec |
| PCIe gen6 (2021) x16 channel | 126.03 GB/sec |
| NVMe Flash bandwidth | ~32 GB/sec |
| GbE link bandwidth | 10 – 100 Gbps |
| | |
| HDD (6Gb/s SAS) | 230 MB/s |
| SSD (6Gb/s SAS) | 550 MB/s |
| 6 Gb/s SAS x4 | 2200 MB/s |
| 12 Gb/s SAS x4 | 4400 MB/s |

Useful Latency Numbers

Initial list from Jeff Dean, Google

| | |
|------------------------------------|----------------|
| L1 cache reference | 0.5 ns |
| Branch mispredict | 5 ns |
| L3 cache reference | 20 ns |
| Mutex lock/unlock | 25 ns |
| Main memory reference | 100 ns |
| Compress 1K bytes with Snappy | 3,000 ns |
| Send 2K bytes over 10Ge | 2,000 ns |
| Read 1 MB sequentially from memory | 100,000 ns |
| Read 4KB from NVMe Flash | 50,000 ns |
| Round trip within same datacenter | 500,000 ns |
| Disk seek | 10,000,000 ns |
| Read 1 MB sequentially from disk | 20,000,000 ns |
| Send packet CA → Europe → CA | 150,000,000 ns |

Total Cost of Ownership (TCO)

TCO = capital (CapEx) + operational (OpEx) expenses

Operators perspective

CapEx: building, generators, A/C, compute/storage/net HW

Including spare parts, amortized over 3 – 15 years


OpEx: electricity (5-7cents/KWh), repairs, people, WAN, insurance, ...

Users perspective

CapEx: cost of long term leases on HW and services

OpeEx: pay per use cost on HW and services, people

Operator's TCO Example



[Source: James Hamilton]

Hardware dominates TCO, so we have to make it cheap
Must use it as well as possible

Reliability

Failure in time (FIT)

Failures per billion hours of operation = $10^9/\text{MTTF}$


Mean time to failure (MTTF)

Time to produce first incorrect output

Mean time to repair (MTTR)

Time to detect and repair a failure

Availability



$$\text{Steady state availability} = \text{MTTF} / (\text{MTTF} + \text{MTTR})$$

Yearly Datacenter Flakiness

- ~0.5 **overheating** (power down most machines in <5 mins, ~1-2 days to recover)
- ~1 **PDU failure** (~500-1000 machines suddenly disappear, ~6 hrs to come back)
- ~1 **rack-move** (plenty of warning, ~500-1000 machines powered down, ~6 hrs)
- ~1 **network rewiring** (rolling ~5% of machines down over 2-day span)
- ~20 **rack failures** (40-80 machines instantly disappear, 1-6 hours to get back)
- ~5 **racks go wonky** (40-80 machines see 50% packet loss)
- ~8 **network maintenances** (4 might cause ~30-minute random connectivity losses)
- ~12 **router reloads** (takes out DNS and external VIPs for a couple minutes)
- ~3 **router failures** (have to immediately pull traffic for an hour)
- ~dozens of minor 30-second blips for dns
- ~1000 **individual machine failures** (2-4% failure rate, machines crash at least twice)
- ~thousands of **hard drive failures** (1-5% of all disks will die)

Add to these SW bugs, config errors, human errors, ...

Key Availability Techniques

| Technique | Performance | Availability |
|--|-------------|--------------|
| Replication | ✓ | ✓ |
| Partitioning (db tables sharding) | ✓ | ✓ |
| Load-balancing | ✓ | |
| Watchdog timers (hang prevention) | | ✓ |
| Integrity checks (check sum etc) | | ✓ |
| Canaries (Canary release) | | ✓ |
| Eventual consistency (inconsistency tolerance) | ✓ | ✓ |

Make apps do something reasonable when not all is right

Better to give users limited functionality than an error page

Aggressive load balancing or request dropping

Better to satisfy 80% of the users rather than none

The CAP Theorem

In distributed systems, choose 2 out of 3

Consistency

Every read returns data from most recent write

Availability

Every request executes & receives a (non-error) response

Partition-tolerance

The system continues to function when network partitions occur (messages dropped or delayed)

Useful Tips

Check for single points of failure

Keep it simple stupid (KISS)

The reason many systems use centralized control

If it's not tested, do not rely on it

Question: how do you test availability techniques with hundreds of loosely coupled services running on thousands of machines?