

# Serverless concept

Serverless is a managed, native cloud architecture that shifts operational responsibility of servers to the cloud. That allows you, the developer, to focus on building applications and services without even thinking about the servers and software needed to run your code.

Cost-efficiency is intertwined with serverless because you only pay for what you use, whereas even virtual servers cost money to run 24/7.

In this lesson, we will use a serverless frontend called **API Gateway** which will send the request to a **Lambda Function**, which will serve as our backend.

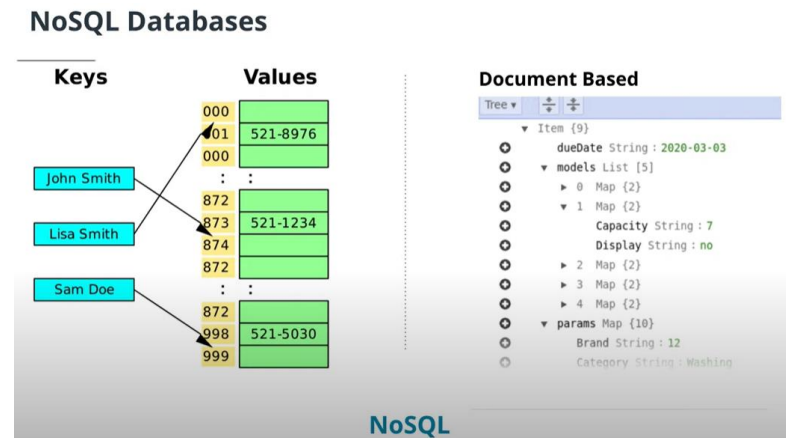
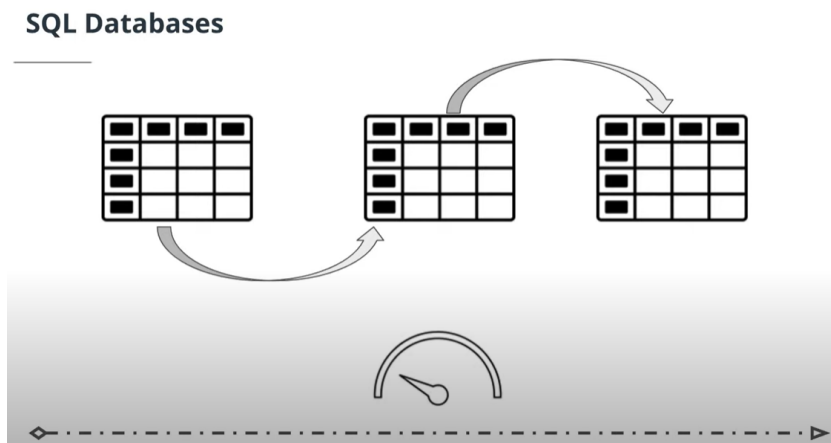
That Lambda will either set or get the user's name from a managed, in-memory database service. The Lambda will return the name back to the user via the API Gateway.



# ElastiCache

## ElastiCache

- Managed, in-memory NoSQL database
- For security, instances are deployed to a private subnet by default and lack a publicly reachable address
- Redis is an open-source data store, whereas ElastiCache is AWS's proprietary service that hosts Redis databases for AWS users
- Great choice for real-time applications such as gaming, chat, and video
- Easily scalable and replicable



Not only SQL databases  
flexible data requirements

# Exercise: ElastiCache For Redis

In this exercise, you will launch a Redis instance on ElastiCache. You will continue to use this instance in the next exercise.

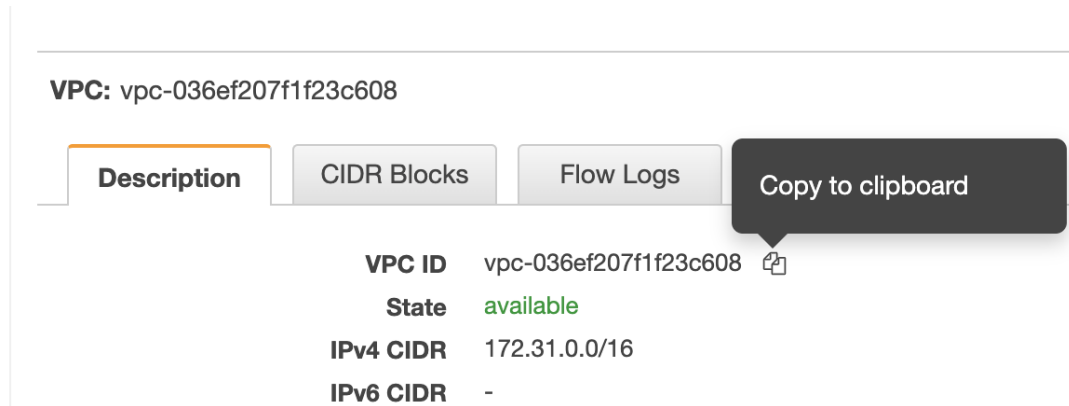
## Instructions

- **Single node** with **no replication**
- **t2.micro** type
- **us-east-1 N.Virginia** region
- Setup on the **default VPC**
- No encryption either **in transit** (No Auth) or **at-rest**
- Name the node **redis**

Before you start, you need to get the Virtual Private Cloud ID of the default VPC:

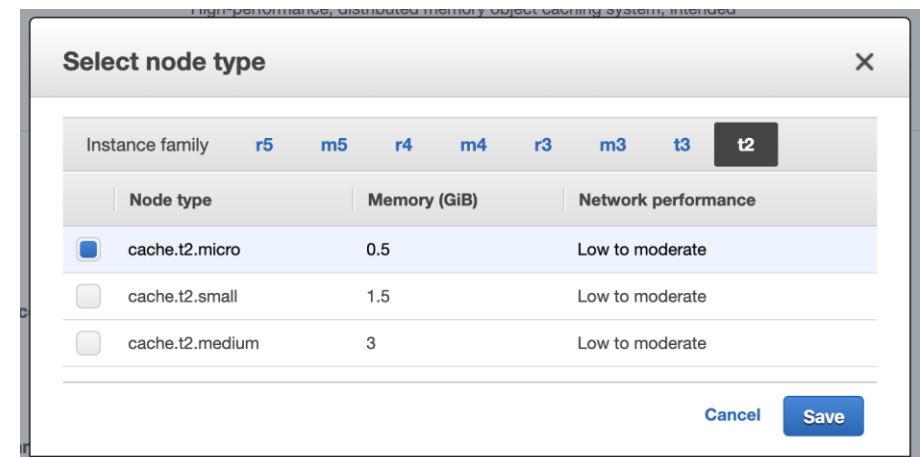
## Default VPC ID

1. Open the [VPC console](#)
2. Note the default VPC ID



## Create the ElastiCache Redis Instance

1. Open the [ElastiCache console](#)
2. Change the region on the top right to be **us-east-1 N.Virginia**
3. Click **Redis** on the sidebar menu, then click the **Create** button
4. Name the instance **redis**
5. Change the number of replicas to 0
6. Change the **Node type** to be **t2.micro**
7. Save



8. Under the Advanced Redis Settings, set the **Name** to be **redis-subnet**
9. Set the **Description** to **Redis Subnet**
10. Change the VPC ID to the default VPC ID that you noted previously
11. Select **us-east-1a** as the subnet

▼ Advanced Redis settings

Advanced settings have common defaults set to give you the fastest way to get started. You can modify these now or after your cluster is created.

Subnet group

Name

Description

VPC ID

Subnets

Subnet ID	Availability zone	CIDR Block
<input type="checkbox"/> subnet-0715be96ca93795d7	us-east-1e	172.31.48.0/20
<input type="checkbox"/> subnet-0281cefdad6db3b7	us-east-1b	172.31.32.0/20
<input checked="" type="checkbox"/> subnet-0e1daeabdb0a4eda1	us-east-1a	172.31.16.0/20
<input type="checkbox"/> subnet-06f592dc54516186a	us-east-1f	172.31.64.0/20
<input type="checkbox"/> subnet-0d1196ce2497a2091	us-east-1c	172.31.0.0/20
<input type="checkbox"/> subnet-09d26c1ef33ce3aca	us-east-1d	172.31.80.0/20

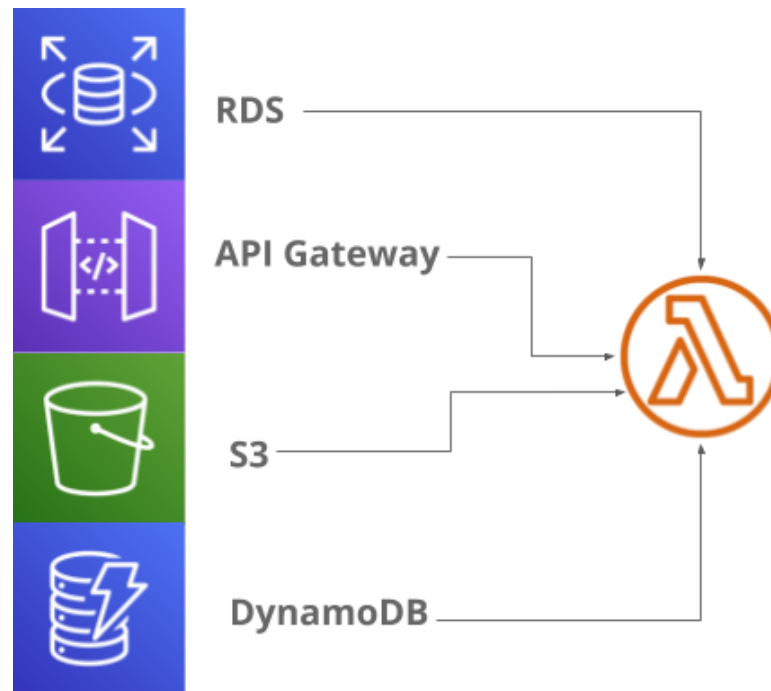
Filter: Search Clusters...										1 to
<input type="checkbox"/>	Cluster Name	Mode	Shards	Nodes	Node Type	Status	Update Action Status	Encryption in-transit	Encryption at-rest	
<input type="checkbox"/>	▼ redis	Redis	0	1 node	cache.t2.micro	available	up to date	No	No	
				Name: redis			Global Datastore: -			
				Global Datastore Role: -			Creation Time: April 9, 2020 at 7:26:23 PM UTC-7			
				Configuration Endpoint: -			Status: available			
				Primary Endpoint: redis.c7icbc.0001.use1.cache.amazonaws.com:6379			Update Status: up to date			
				Engine: Redis			Reader Endpoint: -			
				Engine Version Compatibility: 5.0.6			Node type: cache.t2.micro			
				Availability Zones: us-east-1a			Shards: 0			
				Number of Nodes: 1 node			Multi-AZ: Disabled			
				Description: -			Parameter Group: default.redis5.0 (in-sync)			
				Subnet Group: redis-subnet			Security Group(s): sg-060e56ef0039b4739 (VPC) (active)			
				Notification ARN: Disabled			Maintenance Window: wed:07:00-wed:08:00			
				Backup Retention Period: Disabled			Backup Window: Disabled			
				Encryption in-transit: No			Redis AUTH: No			
				Encryption at-rest: No			Customer Managed CMK: -			

12. Scroll down to the **Backup** section and deselect **Enable Automatic backups**
13. Click on the **Create** button

Wait a few minutes for the instance to initialize. You can refresh the dashboard using the refresh arrows icon in the top-right corner of the dashboard.

# Lambda Function

- Serverless fast solution to run code for a task
- Write code directly in the AWS console or upload a zip package with dependencies
- Supports multiple code languages
- Can be tested from the console



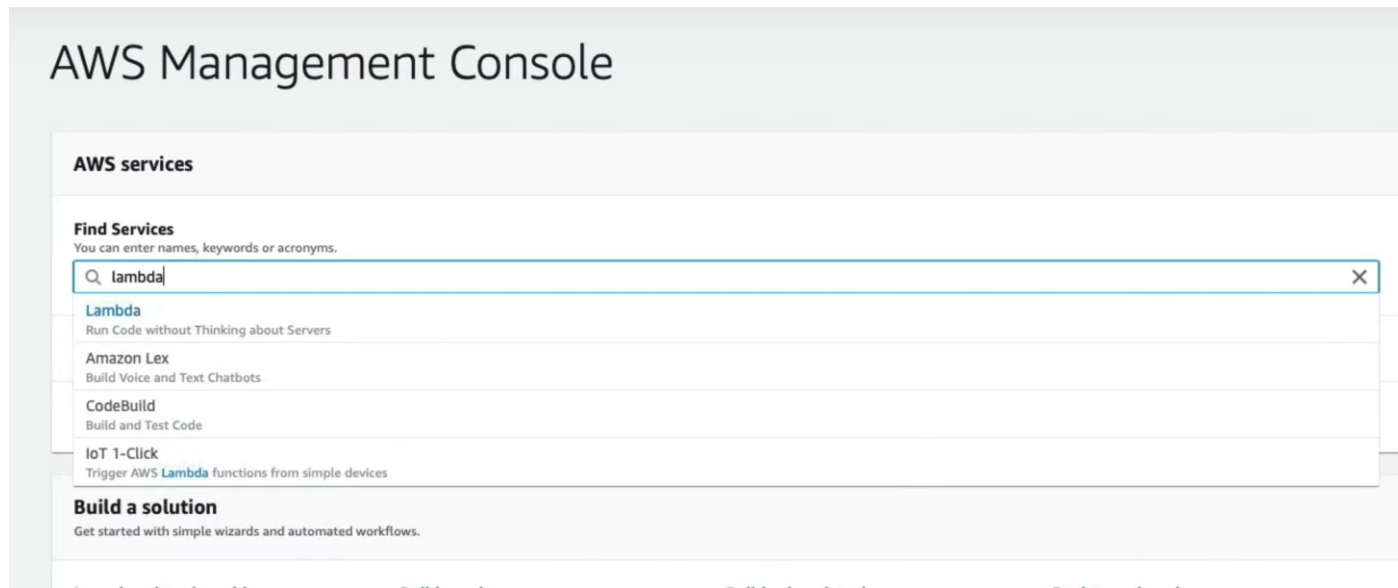
# Exercise: Hello from Lambda

In this exercise, you will create a basic Lambda Function.

## Instructions

- Create a Lambda Function from scratch
- The Lambda Handler should be set to the default **lambda\_function.lambda\_handler**
- Manually trigger the Lambda via a test event and verify the message "Hello from Lambda"

1. Navigate to the Lambda Console by searching for **Lambda** under **Find Services**



2. Click the **Create function** button and select **Author from scratch**
3. Name the function **hello**
4. Set the runtime to **Python3.8**
5. Click the **Create Function** button


Lambda > Functions > Create function

## Create function [Info](#)

Choose one of the following options to create your function.


**Author from scratch** ☒

Start with a simple Hello World example.




**Use a blueprint** ☐

Build a Lambda application from sample code and configuration presets for common use cases.



**Browse serverless app repository** ☐

Deploy a sample Lambda application from the AWS Serverless Application Repository.



### Basic information

**Function name**  
Enter a name that describes the purpose of your function.

hello

Use only letters, numbers, hyphens, or underscores with no spaces.

**Runtime** [Info](#)  
Choose the language to use to write your function.

Python 3.8

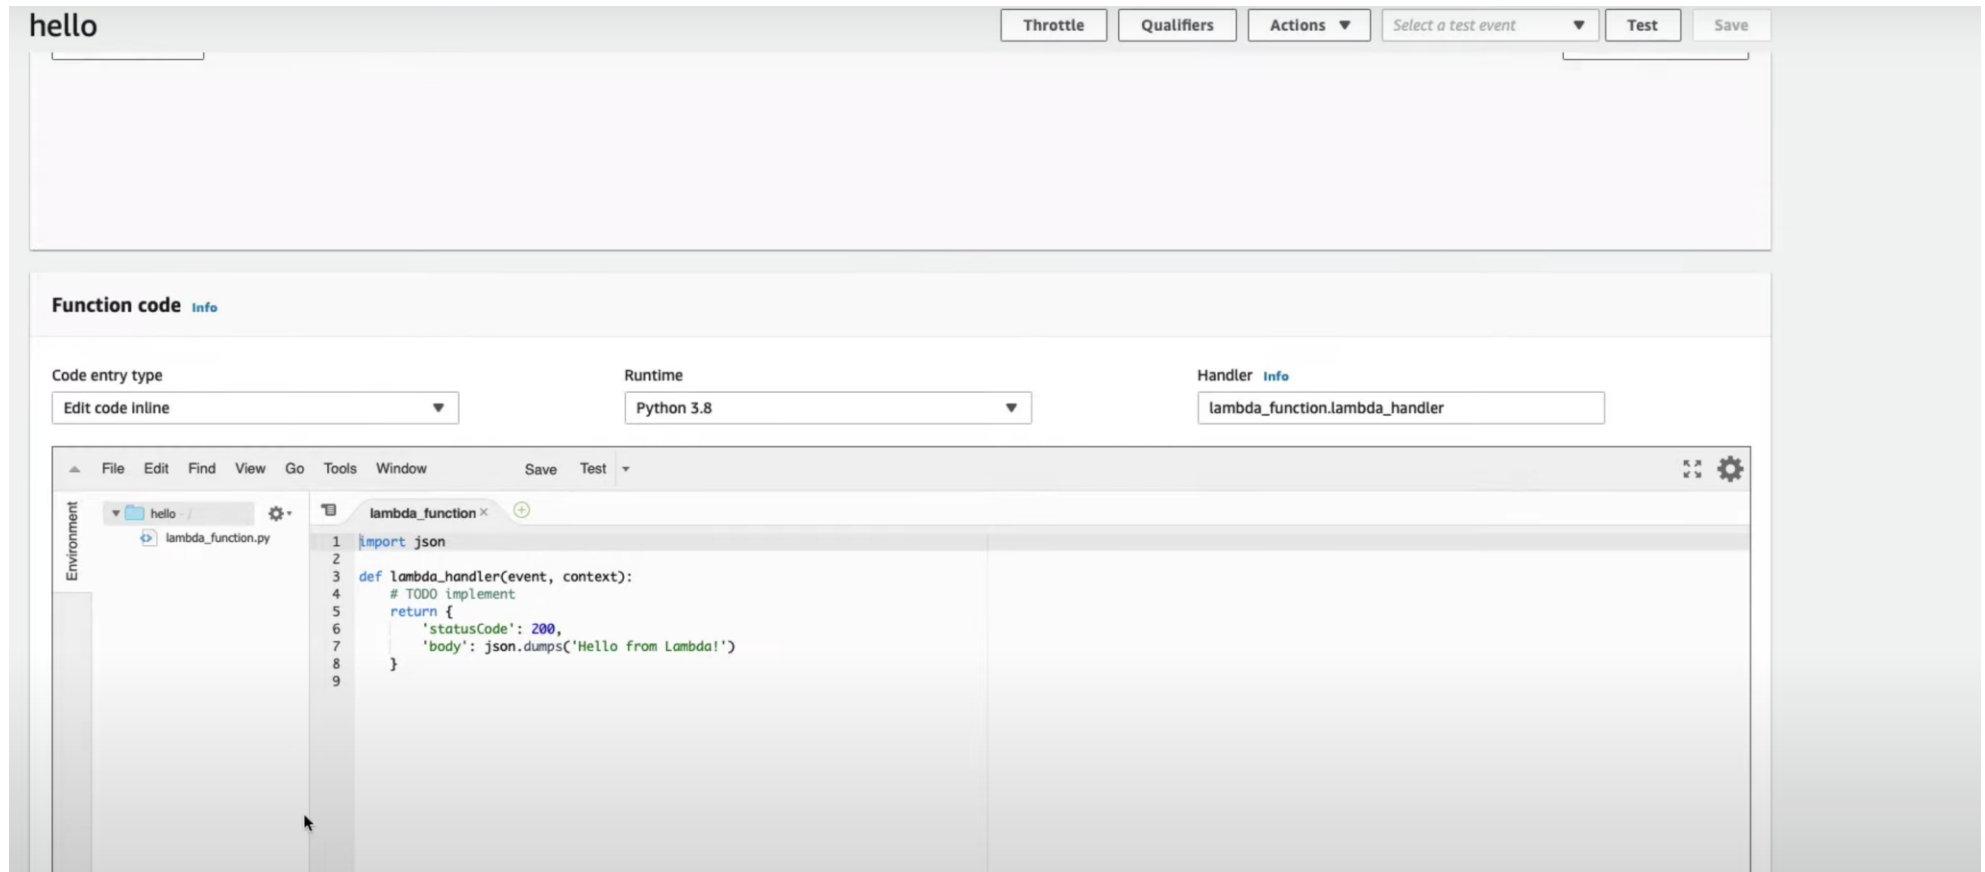
**Permissions** [Info](#)  
Lambda will create an execution role with permission to upload logs to Amazon CloudWatch Logs. You can configure and modify permissions further when you add triggers.

► Choose or create an execution role

Cancel Create function

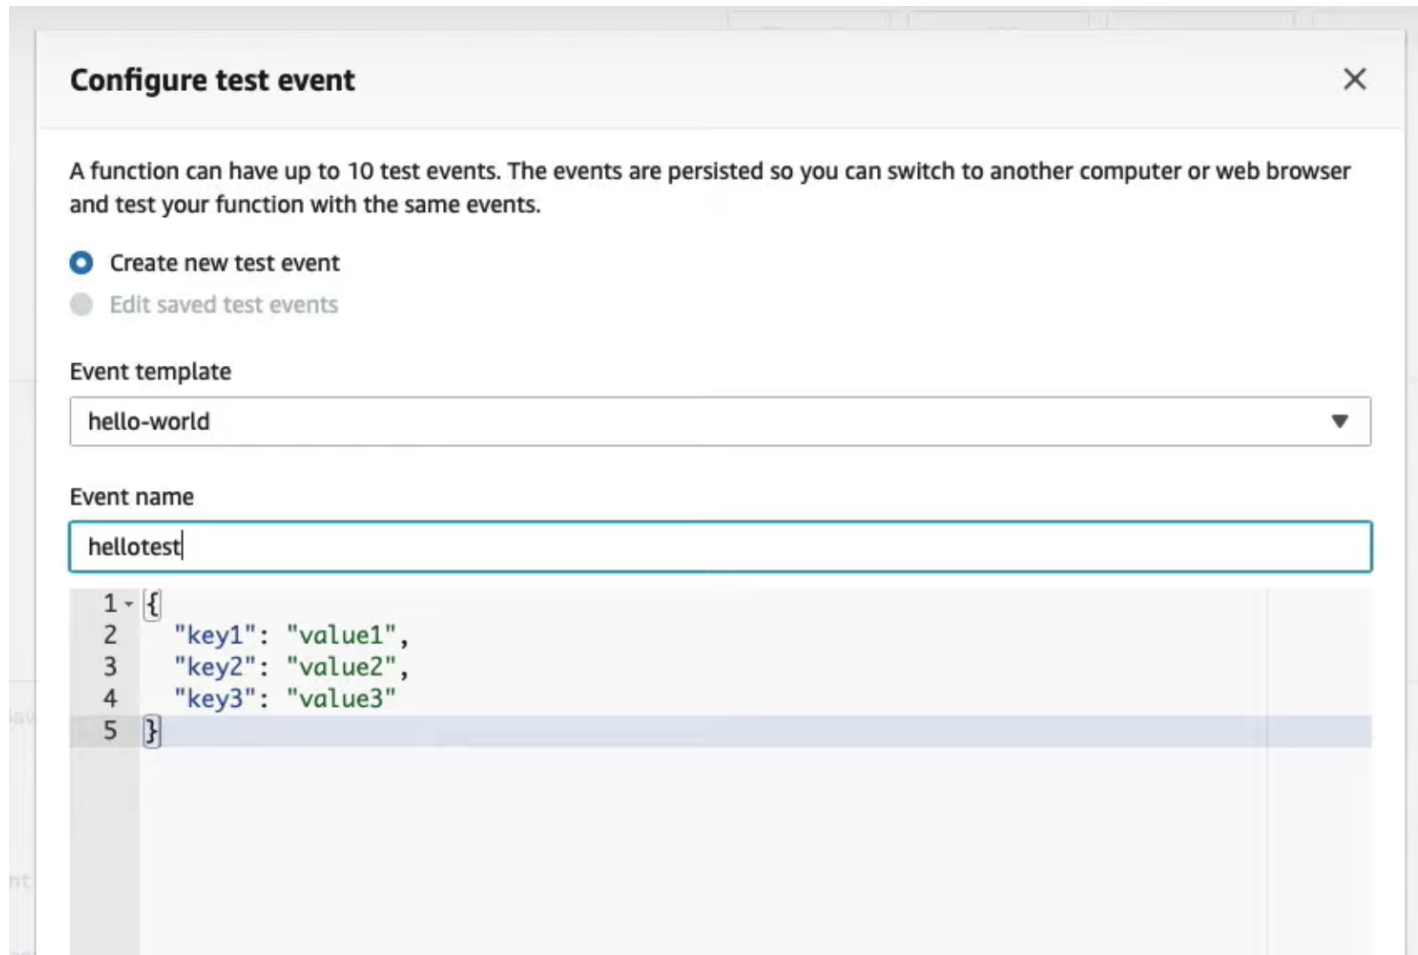


6. Notice the **handler** field on the embedded code page. It's format is **file-name.function-name**. In our case, the file name is **lambda\_function.py** and the code function name is **lambda\_handler**, which is why the handler is pre-set to **lambda\_function.lambda\_handler**.



7. Click on the **Test** button and configure a test event.

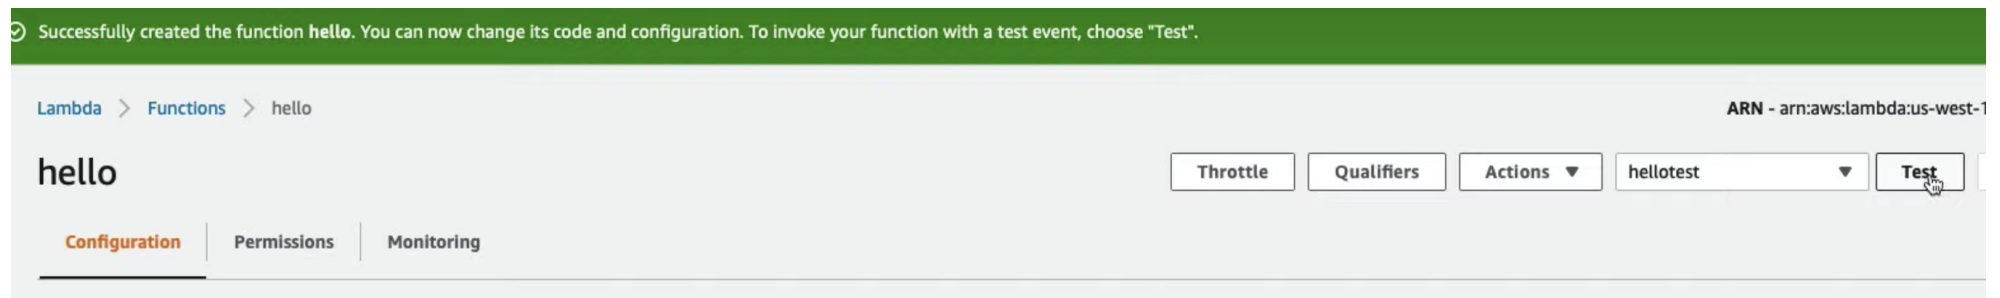
8. Name the event **hellotest** and save the event.



The screenshot shows the 'Configure test event' dialog box. At the top, it says 'Configure test event' with a close button. Below that, a message states: 'A function can have up to 10 test events. The events are persisted so you can switch to another computer or web browser and test your function with the same events.' There are two radio buttons: 'Create new test event' (selected) and 'Edit saved test events'. Under 'Event template', a dropdown menu shows 'hello-world'. The 'Event name' field contains 'hellotest'. Below this is a code editor showing a JSON template:

```
1 {  
2   "key1": "value1",  
3   "key2": "value2",  
4   "key3": "value3",  
5 }
```

9. Click the **Test** button again while the **hellotest** event is pre-selected



The screenshot shows the AWS Lambda console for a function named 'hello'. A green banner at the top says: 'Successfully created the function hello. You can now change its code and configuration. To invoke your function with a test event, choose "Test".' Below the banner, the breadcrumb is 'Lambda > Functions > hello'. The function name 'hello' is displayed. On the right, there are buttons for 'Throttle', 'Qualifiers', 'Actions', and a dropdown menu showing 'hellotest'. A 'Test' button is also visible. At the bottom, there are tabs for 'Configuration' (selected), 'Permissions', and 'Monitoring'.

10. Expand the **Execution result**

11. Verity the message: **"Hello From Lambda!"**

Lambda > Functions > hello ARN - arn:aws:lambda:us-west-1:91546476855

hello Throttle Qualifiers Actions ▼ hellotest ▼ Test Save

✓ Execution result: succeeded (logs) ✕

▼ Details

The area below shows the result returned by your function execution. [Learn more](#) about returning results from your function.

```
{
  "statusCode": 200,
  "body": "\"Hello from Lambda!\""
}
```

**Summary**

Code SHA-256 LD0MjQ0TMan6oFyOz3Y3mJqsPxm9IEhtimXGaDlb8c=	Request ID 32d0d3cc-806e-4c83-b2c1-3bdbb6e96e97
Duration 1.19 ms	Billed duration 100 ms
Resources configured 128 MB	Max memory used 51 MB

**Log output**

The section below shows the logging calls in your code. These correspond to a single row within the CloudWatch log group corresponding to this Lambda function. [Click here](#) to view the CloudWatch log group.

```
START RequestId: 32d0d3cc-806e-4c83-b2c1-3bdbb6e96e97 Version: $LATEST
END RequestId: 32d0d3cc-806e-4c83-b2c1-3bdbb6e96e97
REPORT RequestId: 32d0d3cc-806e-4c83-b2c1-3bdbb6e96e97  Duration: 1.19 ms    Billed Duration: 100 ms Memory Size: 128 MB    Max Memory Used: 51 MB
```

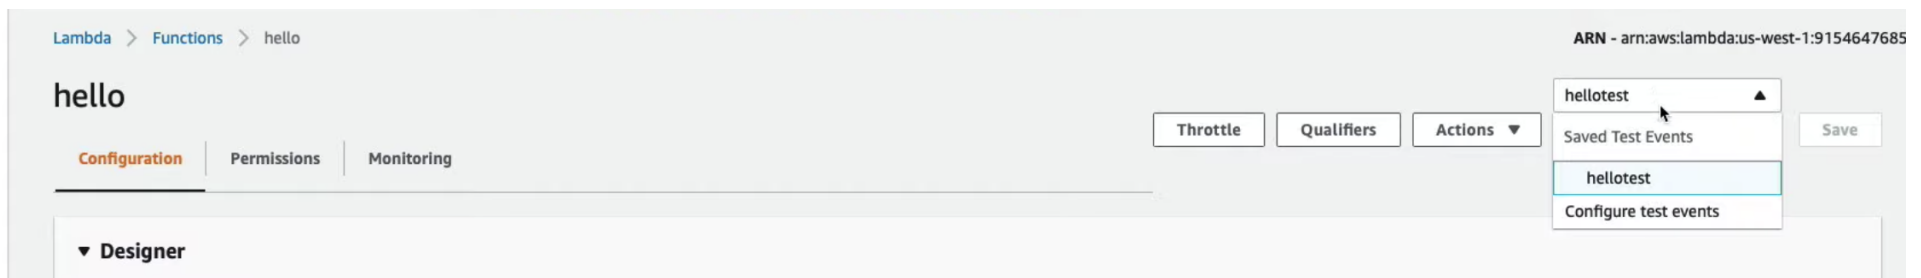
# Exercise: Lambda Events

Modify the **testhello** event for the Lambda with the existing field:

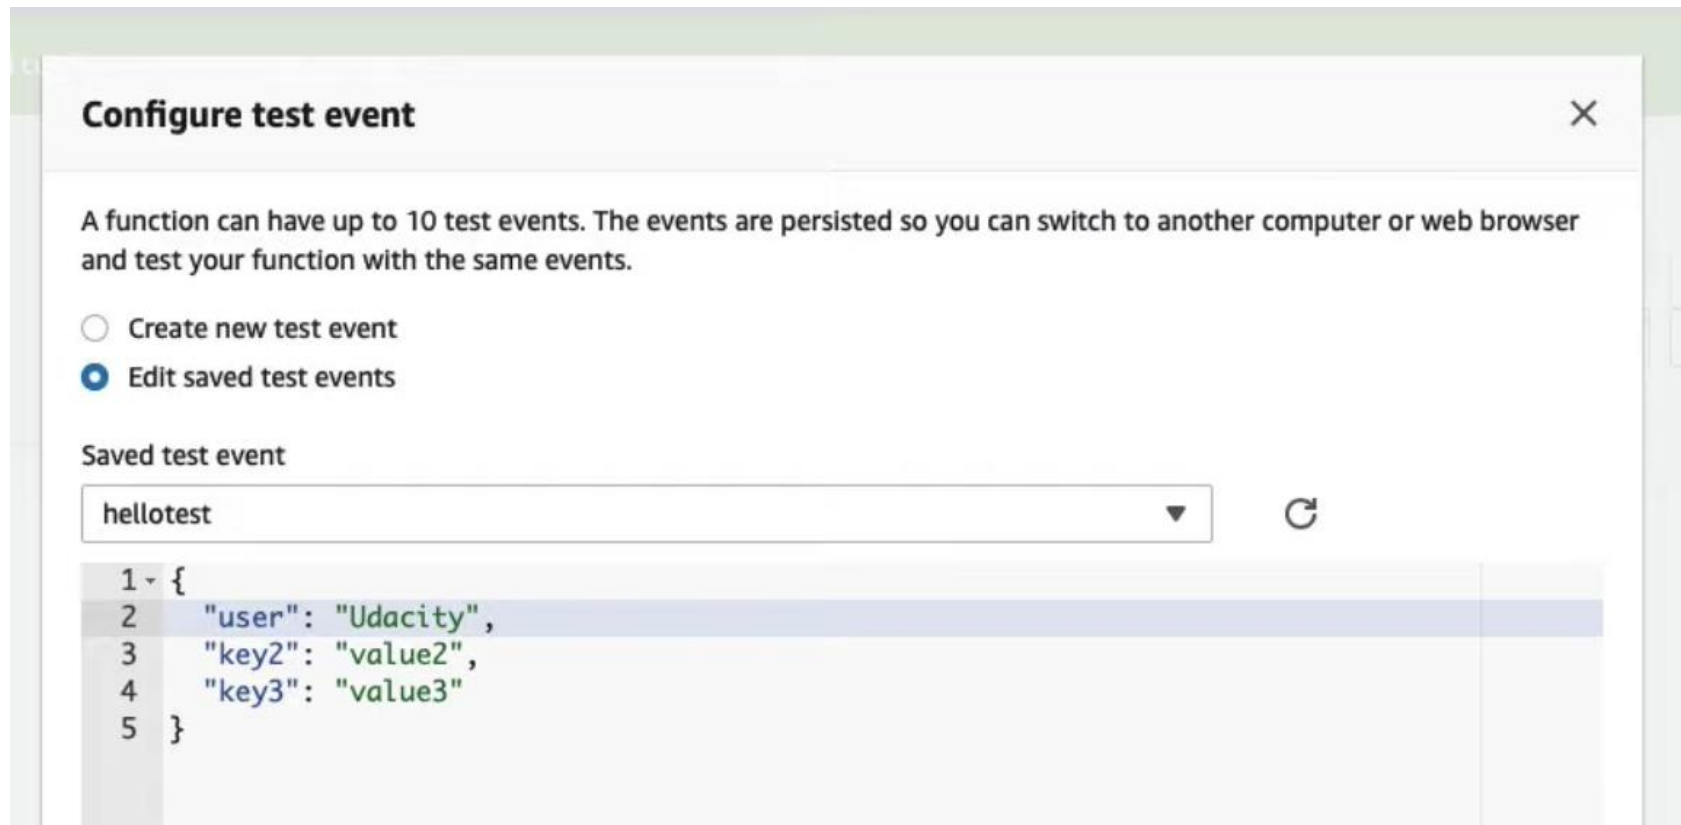
```
{ "user": "Udacity" }
```

- Modify the Lambda Function code to read the user out of the event field and display “Udacity”.

1. From the `hello` Lambda Function, click on the **hellotest** dropdown menu and select **Configure test events**



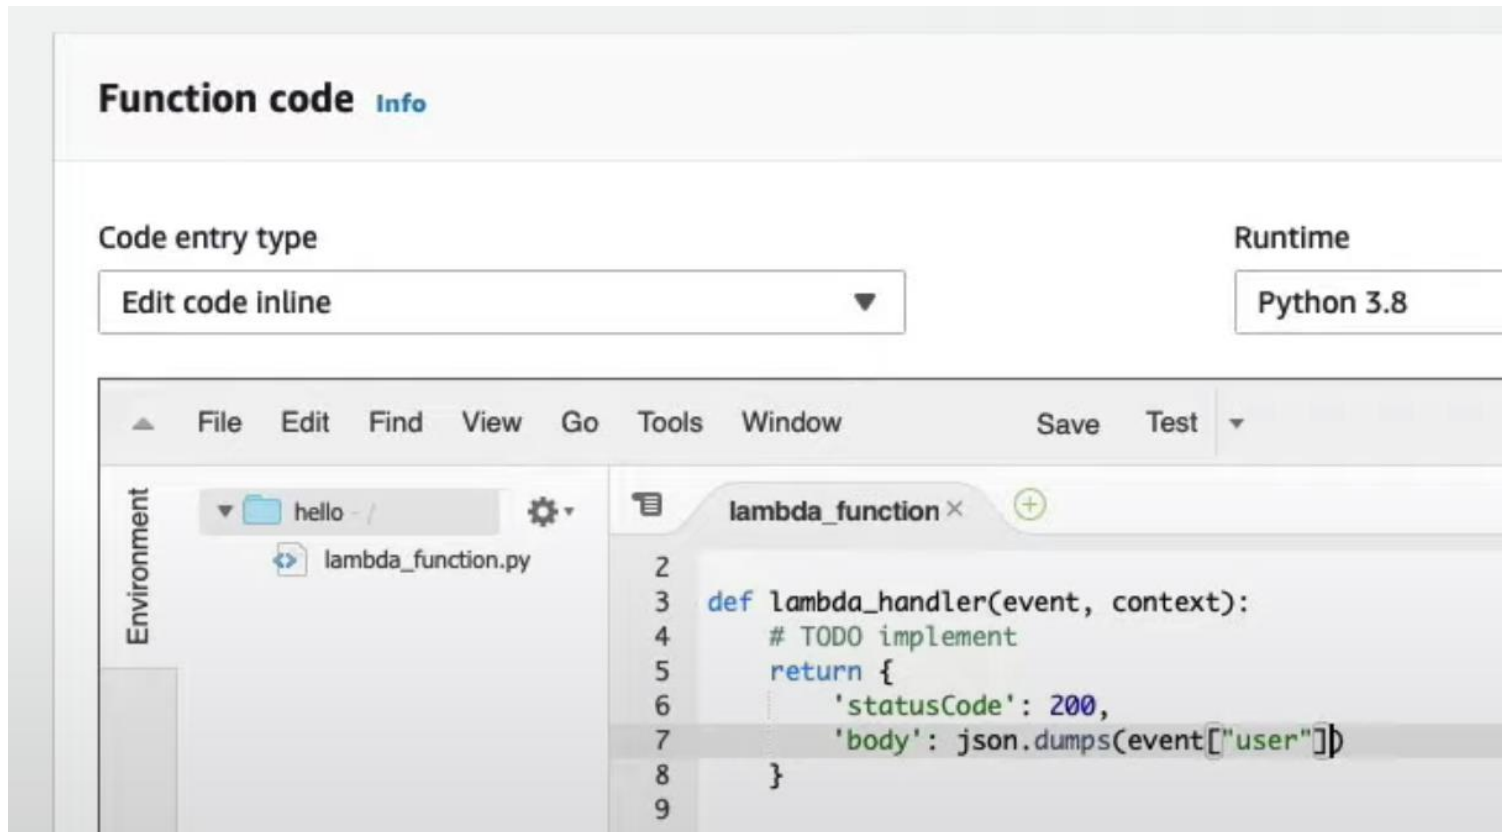
2. Select **Edit saved test events** option (it should be pre-selected already)
3. Change **key1** to be **"user"**
4. Change the value for **"user"** to be **"Udacity"**
5. Click Save Event



6. Edit the code in the editor and set the **body** on **line 7** to be:

```
json.dumps(event["user"])
```

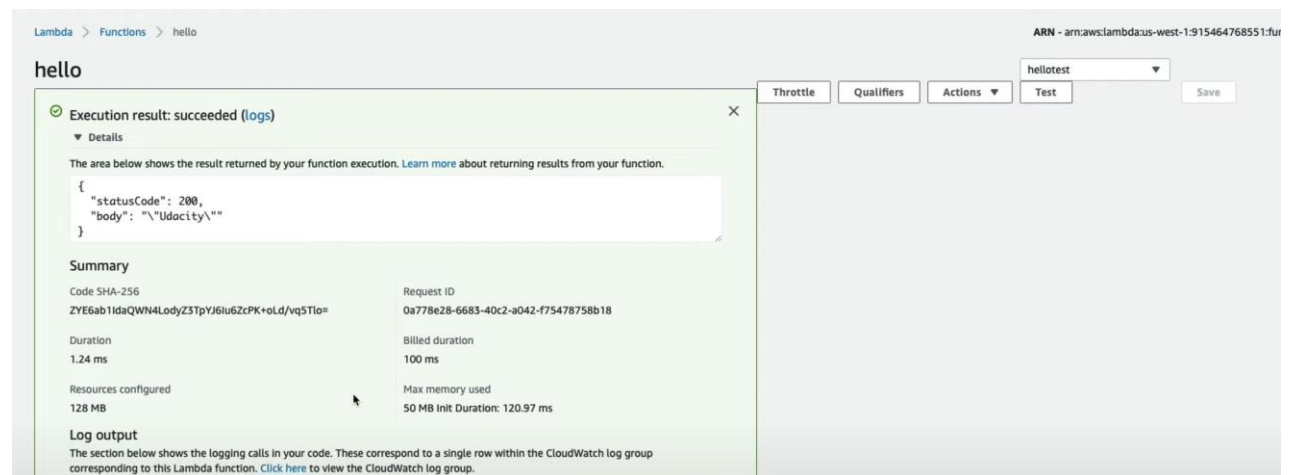
7. Click **Save**.



8. Click the **Test** button.

9. Expand the results.

10. Verify message: **Udacity**



# Exercise: Packaging Code For Lambda Deployment

In this exercise, you will prepare a package for a Lambda deployment. Your code will include both the Lambda Function code *and* the **Redis** dependency.

## Prerequisites

To complete this exercise, you will need:

- [Python3](#)
- [pip](#)

## Instructions

- Install the **Redis** dependency into the **package** folder
  - Hint: use `pip install --target`
- Zip the folder

# Code to include as main.py

```
import redis
import os
import json

def handler(event, context):
    print("Received event: " + json.dumps(event, indent=2))
    redis_host = os.environ.get("REDIS_HOST")
    redis_port = 6379
    redis_password = ""

    r = redis.StrictRedis(
        host=redis_host,
        port=redis_port,
        password=redis_password,
        decode_responses=True
    )

    name = event.get("name")

    if event.get("body"):
        name = json.loads(event["body"]).get("name")

    if name:
        redis_successful_set = r.set("name", name)
        if redis_successful_set:
            return {
                "statusCode": 200,
                "body": "Success! {name} was written to Redis".format(name=name.capitalize())
            }
        else:
            return {
                "statusCode": 500,
                "body": "Oops! Could not write {name} to Redis".format(name=name.capitalize())
            }

    return {
        "statusCode": 200,
        "body": "Hello {name} nice to meet you".format(name=r.get("name").capitalize())
    }
```



## Prerequisites

To complete this exercise solution, you will need:

- [Python3](#) installed

## Create the Lambda Package

For this section, you will need to have Python3 installed.

Lambda Functions are triggered by **events**. When our Lambda Function is triggered, our function will check if the event has a "name" field. if the even has a "name" field, our function will store the name in Redis. If the event does not have a name field, our function will fetch the name set in Redis.

1. Create a new folder on your computer called **hello**
2. Create a new file called **main.py** inside the **hello** folder
3. Copy the code into the **main.py** file (see code on previous slide).
4. Use the following command to install the code dependencies into a folder called **package**  
**pip3 install --target ./package Redis**

5. Create a ZIP archive of the dependencies using the following command

```
~hello$ cd package
```

```
~hello/package$ zip -r9 ../function.zip .
```

6. Add your function code to the archive

```
~/hello/package$ cd ..
```

```
~/hello$ zip -g function.zip main.py
```